

Citadel

Compiling Enterprise Knowledge

Masum Billal

July 7, 2026

Databases Organized Structured Data

The first abstraction layer of enterprise computing was structured data.

Relational databases

SQL | Transactions | Deterministic computation

Applications could compute directly over structured representations.

STRUCTURED DATA BECAME COMPUTABLE BY DEFAULT.

Search Indexed Documents

The second layer introduced unstructured knowledge.

Search engines

Indexing | Ranking | Retrieval

But documents themselves were never made computable.

They were only made searchable.

SEARCH IS NOT A COMPUTATIONAL MODEL. IT IS AN ACCESS LAYER.

AI Still Rebuilds Enterprise Knowledge Every Time

Modern AI systems operate on raw, unstructured documents.

Every system independently reconstructs structure:

OCR

Parsing

Chunking

Embedding

Retrieval

Ranking

This is repeated across every application:

Legal AI, search, copilots, analytics, compliance tools.

NO SHARED REPRESENTATION EXISTS ACROSS APPLICATIONS.

The Missing Layer

Enterprise systems have:

Storage

Search

LLMs

But they lack a computational representation of knowledge.

Core Gap

Documents are not data structures.

They are opaque inputs repeatedly re-interpreted at runtime.

THERE IS NO CANONICAL EXECUTION MODEL FOR ENTERPRISE KNOWLEDGE.

Knowledge Should Be Compiled

Software is not executed directly from source code.



Enterprise knowledge should work the same way.

Compile once. Compute forever.

THE ABSENCE OF A COMPILER IS THE CORE LIMITATION OF ENTERPRISE AI.

Documents Are Syntax

Documents are not data structures.

They are **syntax**.

PDF | DOCX | HTML | EPUB | XML | CSV | Excel | Markdown

Each format is just a different source language.

A compiler should not care about the syntax.

ALL ENTERPRISE DOCUMENT FORMATS ARE JUST SURFACE REPRESENTATIONS.

Canonical Knowledge Representation

Citadel compiles all documents into a single computational form.




Key Idea

Every computation operates on this canonical tree not raw documents.

THIS IS THE ENTERPRISE EQUIVALENT OF INTERMEDIATE REPRESENTATION (IR).

Citadel (The Compiler)

Citadel compiles enterprise knowledge into a canonical representation.



Lossless

- Structure preserved
- Tables preserved
- Provenance preserved

Canonical

- Single representation
- Cross-format unification
- Shared substrate

Computable

- SQL-ready
- API-ready
- Deterministic execution

DOCUMENTS BECOME A COMPUTATIONAL SUBSTRATE.

Knowledge Is Not Retrieved It Is Computed

Traditional systems treat documents as search targets.

Citadel treats knowledge as executable structure.

Query Planning over Canonical Data

THE SYSTEM DOES NOT RETURN TEXT. IT EXECUTES STRUCTURE.

From Retrieval to Knowledge Planning



Not retrieval. **Planning over compiled knowledge.**

SQL Becomes the Execution Layer

All structured knowledge converges into relational execution.

LLM generates queries. Citadel executes them.

Deterministic computation replaces approximation

Tables are first-class citizens

Provenance is preserved end-to-end

MODELS REASON. SYSTEMS COMPUTE.

Wide Net Retrieval Is Not Search

Citadel does not retrieve top-k documents.

It constructs a **bounded evidence field**.

- Query expansion

- Multi-path retrieval

- Structural matching

- Ranking over canonical leaves

IF THE EVIDENCE EXISTS, CITADEL FINDS IT.

Deterministic Knowledge Runtime

Every answer is computed, not guessed.

Canonical IR SQL Execution Answer

No hallucinated structure

No probabilistic aggregation

No re-parsing documents

THE DOCUMENT IS NEVER RE-INTERPRETED.

One Layer. Many Applications.

All enterprise workflows converge onto the same substrate.

Search | Legal AI | Compliance | Analytics | Agents

Key Insight

Applications differ only in logic.
Not in infrastructure.

INFRASTRUCTURE BECOMES SHARED, NOT REBUILT.

One Canonical Layer. Thousands of Applications.

Once enterprise knowledge is compiled,
every application becomes a consumer of the same substrate.

Search | Legal AI | Compliance | Analytics | Agents | ERP

Shift

Applications stop rebuilding infrastructure.
They begin sharing it.

THE CANONICAL LAYER BECOMES THE DEFAULT COMPUTE SURFACE.

The Economic Shift

Today, every AI product rebuilds the same infrastructure.

OCR | Parsing | Retrieval | Ranking | Embeddings | Citations

Each system independently pays for the same primitives.

This is duplicated computation at global scale.

FRAGMENTATION IS THE DEFAULT COST STRUCTURE OF CURRENT AI.

With Citadel

Infrastructure is compiled once.

- One ingestion system

- One canonical representation

- One execution runtime

Applications become thin logic layers:

- Workflows

- UX

- Business logic

- APIs

THE VALUE SHIFTS FROM INFRASTRUCTURE TO EXPRESSION.

The Network Effect of Knowledge

Each document improves the entire system.

Documents Canonical Layer Applications More Documents

Every ingestion improves structure inference

Every query improves planning quality

Every application benefits from shared representation

THE DATASET IS NOT STATIC IT IS A COMPILING SYSTEM.

Why This Becomes a Moat

Citadel improves with scale in ways competitors cannot replicate.

Structural inference improves with corpus diversity

Table extraction improves with schema variety

Query planning improves with workload history

Provenance graph strengthens with usage

Core Property

The system becomes more accurate as it compiles more reality.

THIS IS NOT A MODEL MOAT. IT IS A REPRESENTATION MOAT.

Why Incumbents Cannot Easily Rebuild This

Existing systems assume documents are the unit of computation.

Citadel replaces that assumption entirely.

- Search engines index text

- Vector DBs embed chunks

- LLM apps re-interpret documents every query

None of these systems compile structure.

THE SHIFT IS ARCHITECTURAL, NOT INCREMENTAL.

The Fourth Enterprise Data Layer

Enterprise computing has evolved through distinct abstractions:

Relational Databases

Data Warehouses

Search Engines

Citadel: Canonical Knowledge Layer

Definition

A compiled, executable representation of enterprise knowledge.

THIS IS THE MISSING PRIMITIVE IN ENTERPRISE COMPUTING.

Closing

Every major computing era introduced a compiler/runtime:

Programming languages compilers/interpreter/machine code

Data databases/SQL

Web search engines/retrieval

Enterprise knowledge now needs its own compiler/runtime.

Citadel is that compiler and runtime.

NOT SEARCH. NOT RETRIEVAL. A COMPUTING PRIMITIVE.