

Citadel

The Deterministic Knowledge Layer for Enterprise Data

Masum Billal

2026

Enterprise Data is Not Usable at Query Time

Companies already have the data:

- PDFs, Excel models, contracts, reports

But when leaders ask real questions:

What changed in revenue across all reports since 2019?

the answer is still:

- manual reconstruction
- spreadsheet rebuilding
- human interpretation

Enterprise knowledge exists — but it is not computable.

Why Existing AI Systems Fail

Modern AI pipelines:

- OCR → text
- chunking → embeddings
- LLM → reasoning

But one thing is lost permanently:

structure

Tables become text. Layout disappears. Relationships break.

So AI is forced to guess.

Guessing breaks financial, legal, and audit-grade workflows.

The Real Problem

Enterprise workflows today are not search problems.
They are reconstruction problems:

- re-reading documents
- rebuilding tables manually
- re-validating numbers repeatedly

Result:

- days to answer simple questions
- inconsistent outputs across teams
- high operational risk

The bottleneck is reconstruction, not access.

We don't need better document understanding.

We need a new primitive:

Documents should behave like databases

Not text. Not embeddings. Not summaries.

But structured, queryable data.

The Citadel Approach

Citadel converts documents into structured systems of record.

We ingest:

- PDFs
- Excel files
- Word documents
- scans, CSV, JSON

And compile them into:

a relational, queryable enterprise knowledge layer

How It Works

- ① Ingest any document format
- ② Extract full structure (tables, layout, text)
- ③ Store as relational data (Postgres)
- ④ AI generates SQL queries
- ⑤ Database returns exact answers

AI writes queries. Database computes truth.

What This Unlocks

Citadel enables:

- instant cross-document financial analysis
- contract risk discovery at scale
- audit-ready traceable answers
- enterprise-wide data comparison in seconds

What used to take days now takes seconds.

Why Citadel Wins

Most systems:

- improve search
- improve embeddings
- improve summarization

Citadel removes the problem entirely:

We make documents directly queryable as data

No reconstruction. No guessing. No ambiguity.

Why Now

Three shifts:

- AI is now strong at reasoning and query generation
- enterprises are drowning in document-based data
- demand for secure, on-prem systems is rising

The system gap is now large enough to fund a new layer.

Vision

In the future:

- documents are not read
- they are queried
- AI operates on structured enterprise knowledge

Citadel makes enterprise knowledge computable.