

Citadel

The Operating System for Private Organizational Knowledge

Masum Billal

July 1, 2026

Every Company Already Owns the Answers...

- **Trapped in Millions of Files:** Contracts, PDFs, invoices, emails, spreadsheets, scanned documents, CAD drawings, and compliance reports.
- **Yet Nobody Can Ask Concrete Questions:**
 - *"Which suppliers increased prices more than 15% over the last three years?"*
 - *"Find every contract exposing us to GDPR liability."*
- **95% of enterprise knowledge is effectively unsearchable.**

Why Today's AI Still Fails

LLMs & Deep Research

- Require manual file uploads.
- Forget information over long chats.
- Hit strict token/context limits.

Naive RAG

- Text chunking destroys layouts.
- Fails completely on large tables.
- Poor cross-document reasoning.

Enterprise Search & Glean

- Returns links to documents, not verified answers.
- Relies on keyword/simple matching.

Legacy OCR

- Extracts raw text strings but completely loses meaning and relations.

Companies don't need another chatbot. They need an intelligence layer.

Citadel: Turn Unstructured Docs into Structured Intelligence

A search engine for your entire private knowledge base.


Upload millions of documents once. Ask questions forever.

Every single answer is:


- **Traceable:** Labeled back to the exact source coordinate.
- **Structured:** Extracted cleanly into usable data tables.
- **Explainable:** Driven by evidence, not LLM guesses.
- **Private:** Run completely locally or in an isolated sandbox.

The Workflow Shift

Without Citadel (Hours Lost)



With Citadel (Seconds Saved)



Built on the World's Best Open OCR — Made Faster

95.04

Official OmniDocBench Score

Top-3 Worldwide (Beating
GPT-4o, Gemini 3, and Mistral
OCR)

2.2×

Faster Throughput

Than stock MinerU pipeline using
the exact same model and GPU

Production-Ready Infrastructure

Tested over an official **1,651-page enterprise benchmark** with **0 timeouts or processing failures**. Runs entirely locally with zero cloud dependencies required.

Why Citadel is Different

We replaced traditional RAG text-splitting with a three-layer production architecture:

- 1. Intelligent Ingestion:** Natively preserves multi-dimensional tables, layouts, formulas, and visual relationships instead of flattening documents into generic text.
- 2. Structured Knowledge Engine:** Creates searchable, interconnected data objects instead of isolated, fragmented text chunks.
- 3. Evidence-First Retrieval:** Answers are synthesized strictly from verified evidence clusters, preventing standard vector embedding hallucinations.

One Engine. Every Domain.

Finance: *"Which vendors increased their line-item prices after the Q3 inflation spikes?"*

Legal: *"Show me every active counterparty contract that contains an automatic renewal clause."*

Manufacturing: *"Compare all field maintenance reports from the last decade mentioning bearing failures."*

Healthcare: *"Identify all local patient records meeting the precise clinical inclusion criteria."*

HR & Ops: *"Which candidates in our resume database mention both Rust and Kubernetes?"*

Private by Design

- **Zero Training Guardrails:** Customer document assets are never utilized to train, fine-tune, or modify underlying model weights.
- **Bring Your Own LLM:** Connects ephemerally to your existing enterprise infrastructure (OpenAI API, Azure, or entirely local weights).
- **Walled Sandbox Indexing:** Files are processed on-the-fly into localized layout structures with absolute data isolation.
- **Cloud Optional:** Fully operational inside on-premise air-gapped secure facilities or private clouds.

Why This is Difficult to Copy

Instead of generic AI wrappers, our moat is built on years of foundational ingestion engineering:

- **High-Performance Parsing Pipeline:** Custom state-of-the-art layout preservation engine.
- **Deterministic Table Reconstruction:** Keeps structural cell geometry intact without breaking text strings or losing header hierarchies.
- **Evidence Graph Layer:** Maps real-time cross-document relationships before synthesis.
- **Massively Parallel Inference Stream:** Custom streaming implementation optimized to run 128-wide concurrent processing streams.

Scalable Business Model

Tier	Delivery Mechanism	Pricing Metric
Desktop Pro	Native Application	Predictable Sub
Cloud Scale	Managed SaaS Infrastructure	Consumption U
Enterprise	On-Prem / Air-gapped Private Cloud	Per-Seat Licens
API Layer	Developer Infrastructure Pipeline	Metered Per-Pa

Go-To-Market Friction Reduction

Product-led growth starting at the developer and professional desktop level, scaling seamlessly into total corporate data infrastructure.

The Long-Term Vision

Today: Search, aggregate, and instantly extract insights across all your unstructured corporate documents.

Tomorrow: Every scaling enterprise runs a continuous, private, self-updating knowledge engine over their entire historical data footprint.

Eventually: Citadel becomes the foundational, secure infrastructure layer sitting between human knowledge and AI reasoning models.